

# Kubernetes Networking

Learn The Ins And Outs Of Networking In K8s

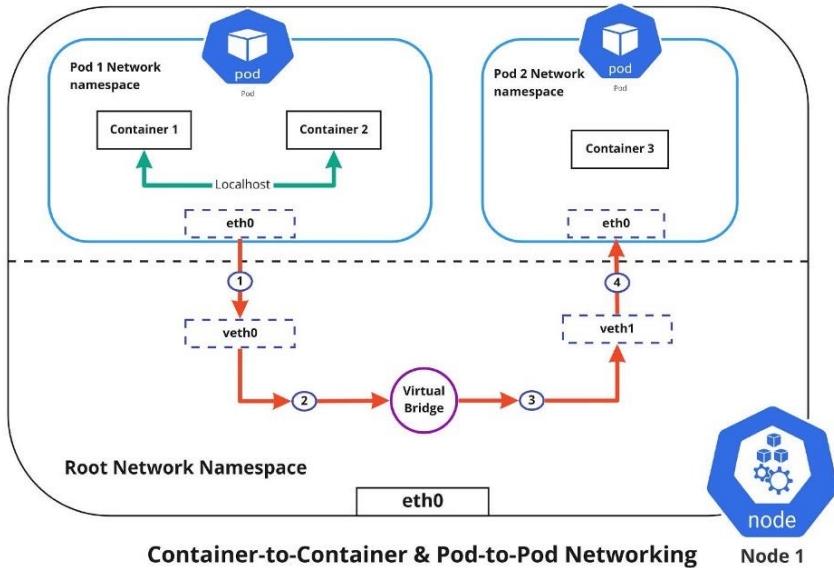


**Michael Levan**

Trainer, Course Creator, and coach

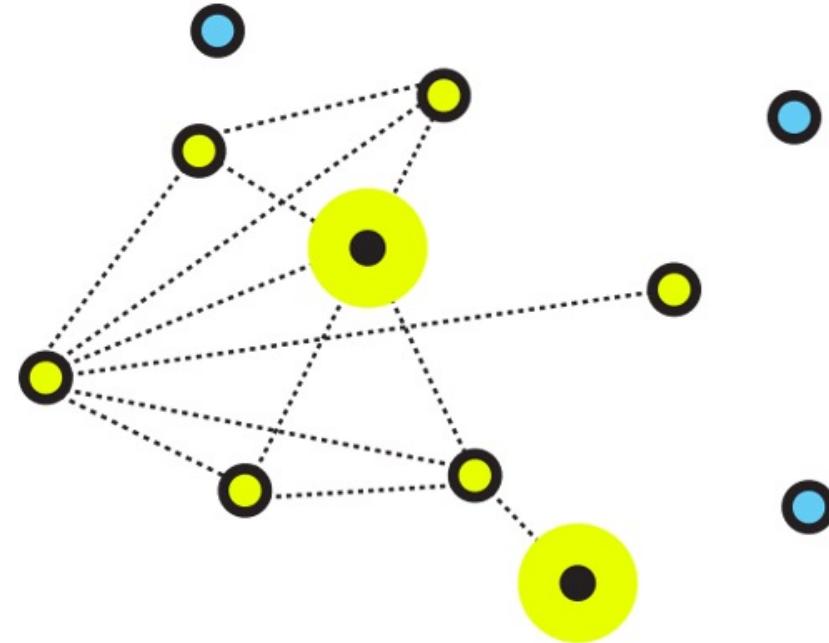
# Section 1: Kubernetes Host Networking

- Cluster Networking
- Cloud Networking
- Static Pods



# Networking Crash Lesson

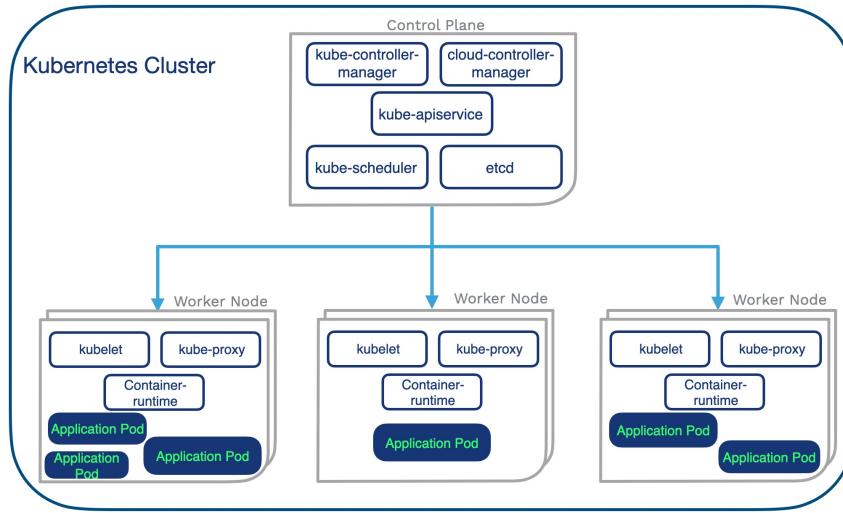
- CIDRs
- IP's
- Subnets
- OSI model



# Cluster Networking

- The cluster itself
- IP addresses and CIDR's for the Worker Nodes
- IP addresses and CIDR's for the Control Planes

# The Cluster Itself



- Control Planes
- Worker Nodes
- Subnets and how it's all connected

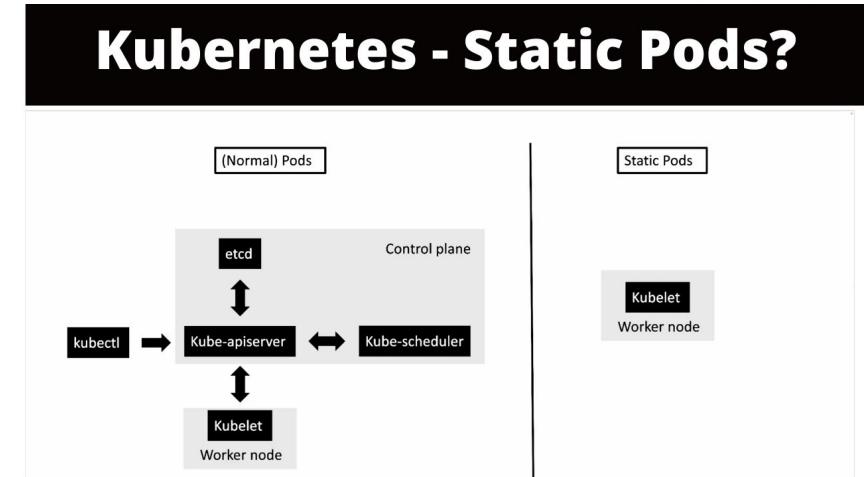
# Quick Note On Kube-Proxy and CoreDNS

- kube-proxy: Runs on every node and is responsible for local cluster networking.
- CoreDNS: The cluster's DNS service has a static IP address that is hard-coded into every Pod on the cluster

# Static Pods When The Cluster Is Created

- Bound to the Control Plane

## Kubernetes - Static Pods?



# Bare Metal vs VM (Cloud Later)

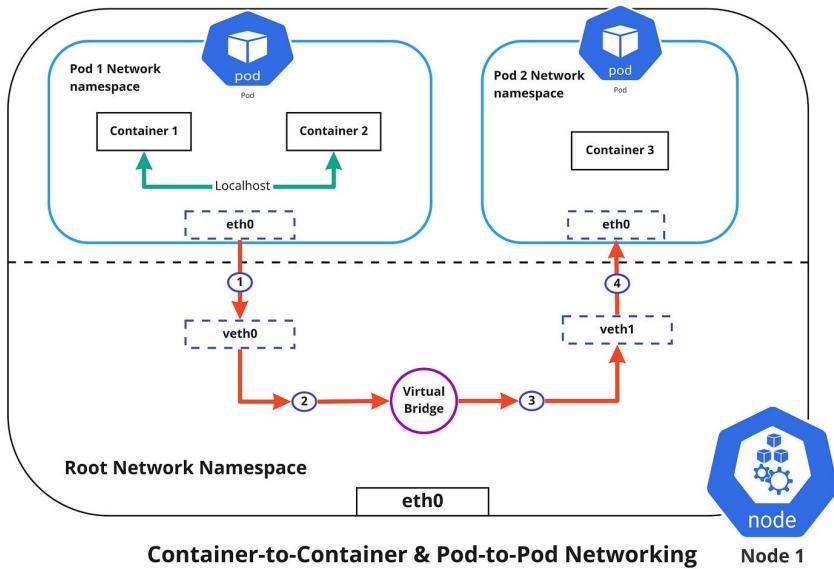
- Everything is managed by you

# Cloud Networking

- Switches, routers, firewalls, etc are all virtual

# Networking With Managed Kubernetes

- Managed by you



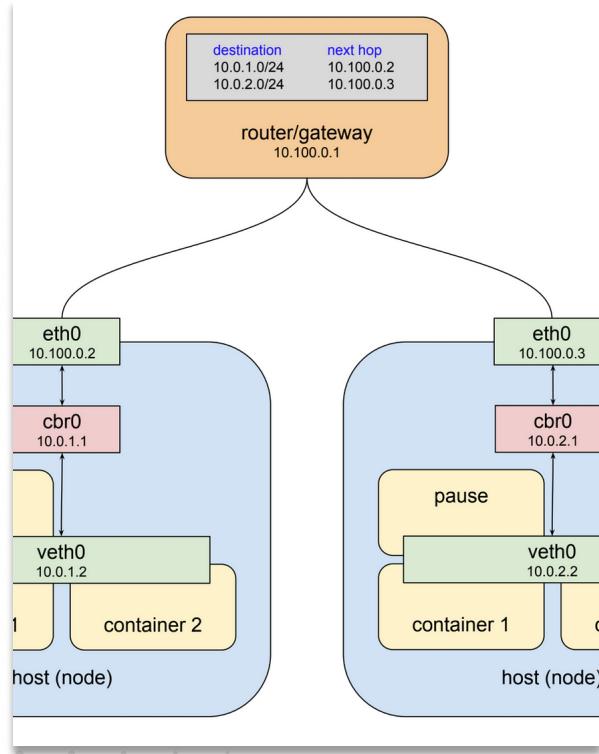
# Secure Cluster Networking

- Proper firewall rules
- Private IPs for the Control Plane

# Lab

- Cluster networking in the cloud (Azure or AWS)

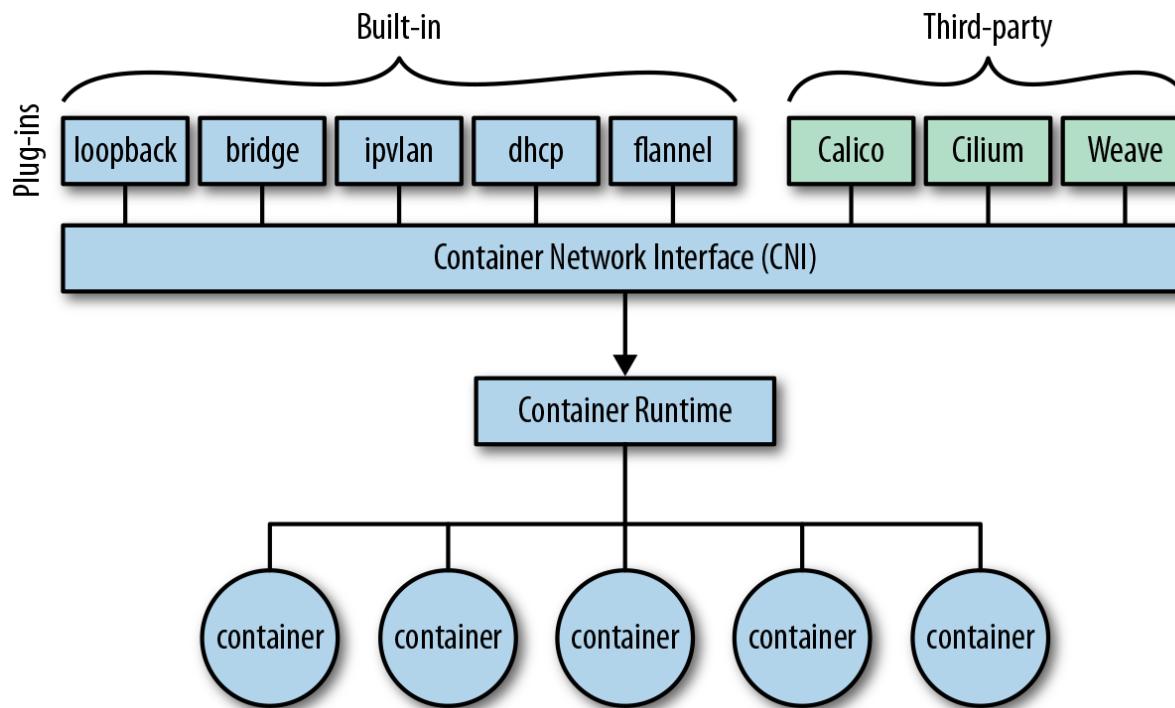
## Section 2: Kubernetes Internal Networking



- CNI
- kube-proxy
- Pod networking
- Pod and Service communication
- Ingress
- eBPF

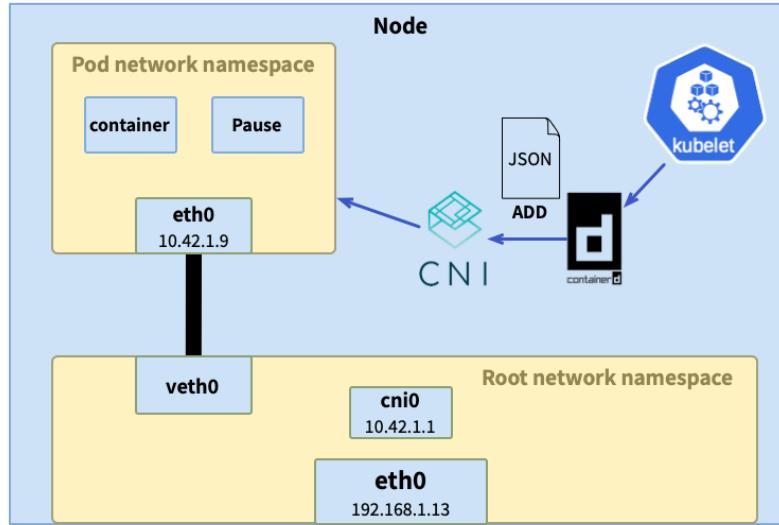


# Container Network Interface (CNI)



# Implementing a CNI

- Options (available CNIs)
- Installation methods



# CoreDNS



# CoreDNS

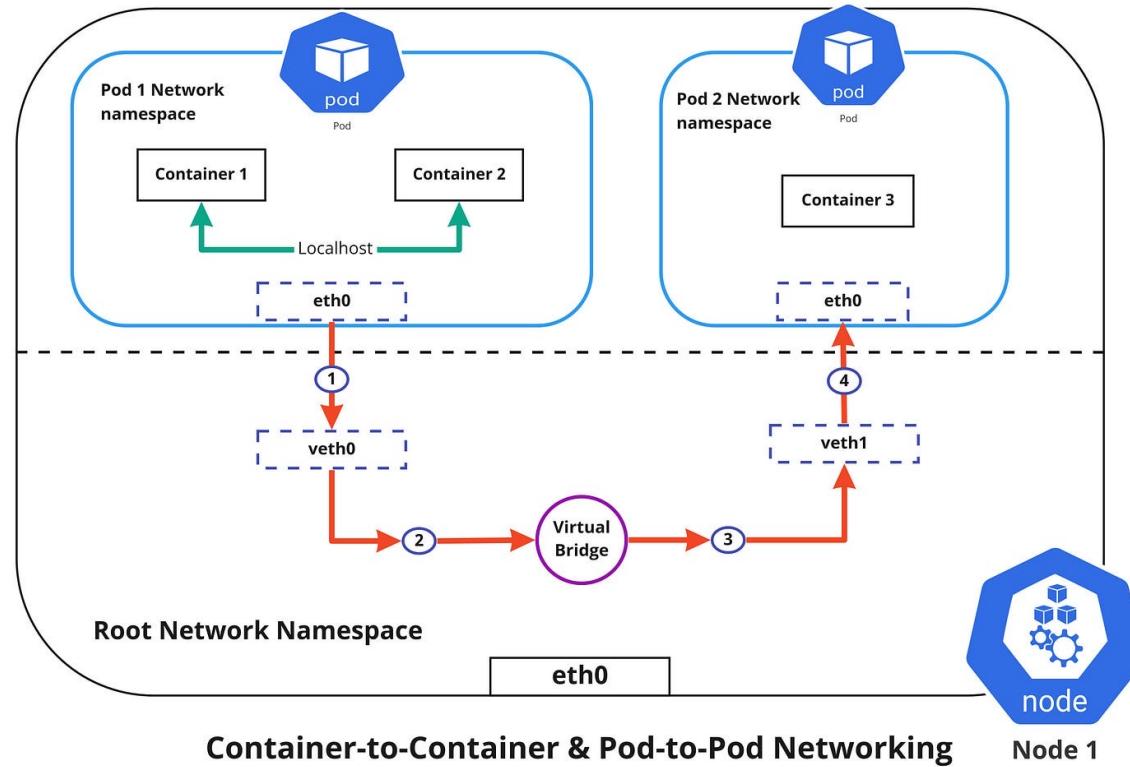
# Pod Networking

- How Pods get IP addresses
- Lease times
- Container IPs
- Sidecar IPs

# Service Networking

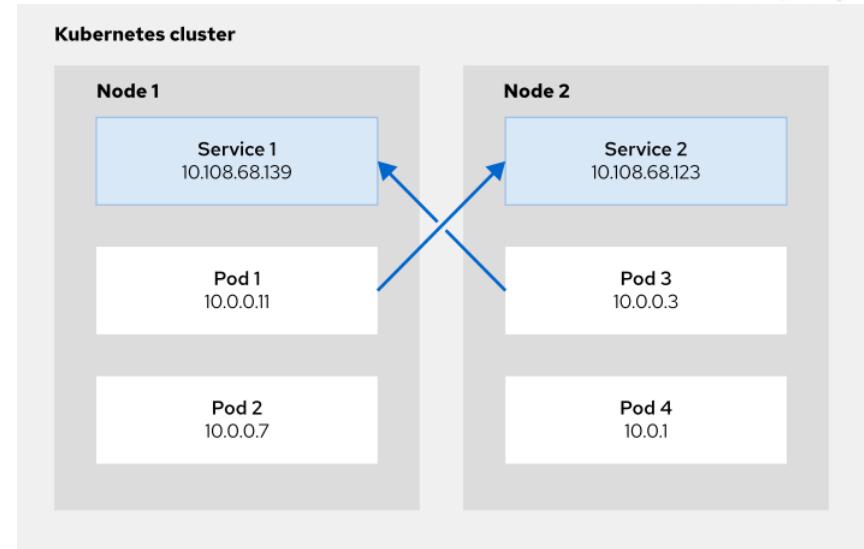
- How services get IPs
- How services get DNS
- How services talk to Pods

# Pod To Pod Communication



# Service To Service Communication

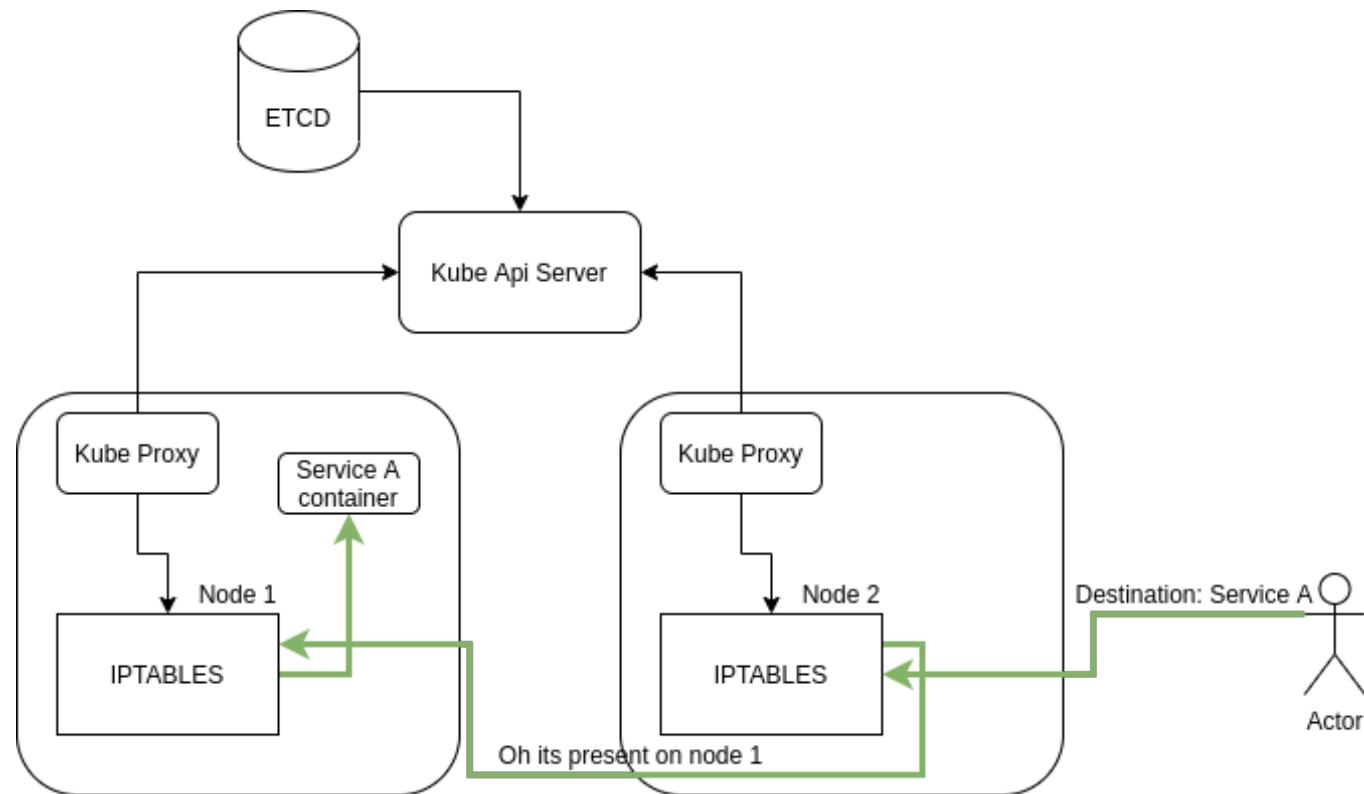
- Via IP or DNS



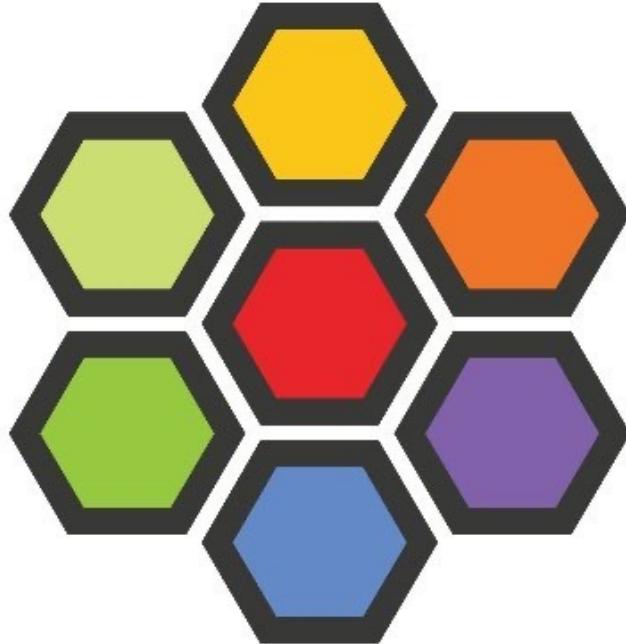
# How To Secure It (It's all Unencrypted)

- Service Mesh for Services
- eBPF for Pods

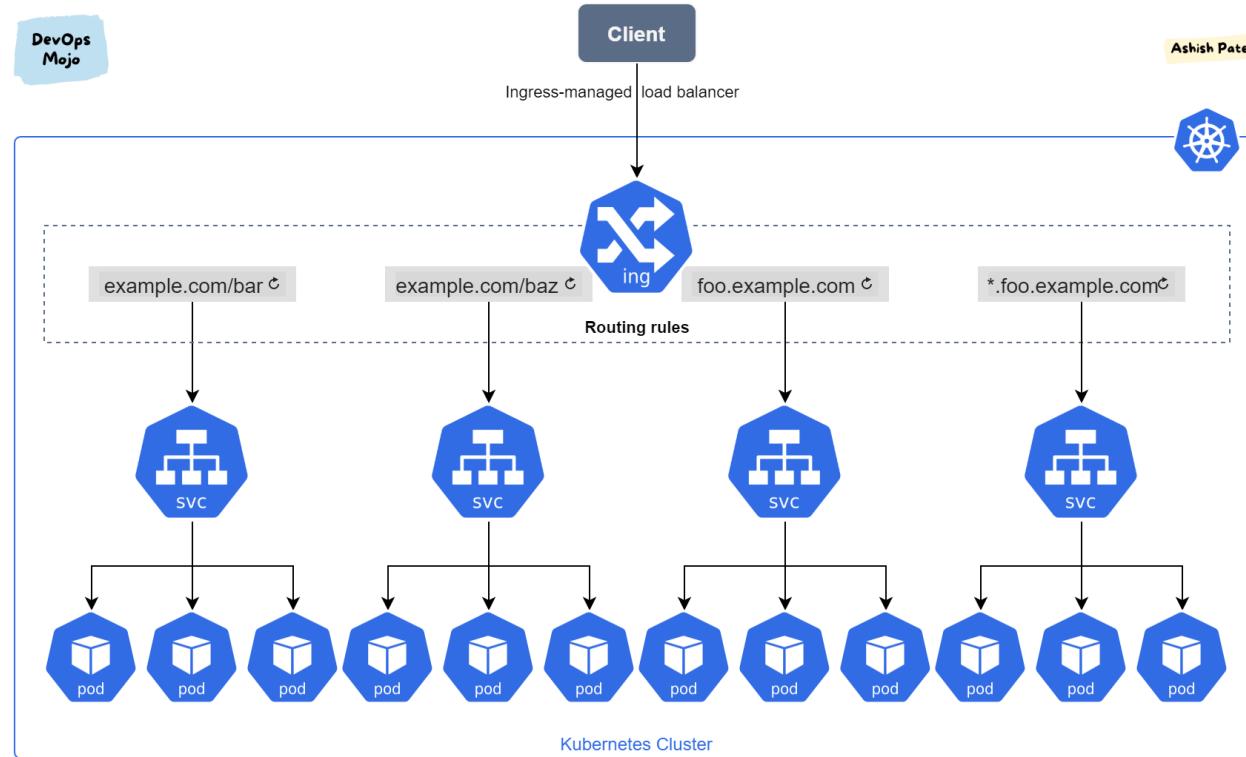
# Kube-Proxy



- Secure Pod communication with a security-centric CNI



# Ingress



# Ingress Implementation

- AKS Application Gateway Ingress Controller is an ingress controller that configures the Azure Application Gateway.
- Alibaba Cloud MSE Ingress is an ingress controller that configures the Alibaba Cloud Native Gateway, which is also the commercial version of [Higress](#).
- Apache APISIX ingress controller is an Apache APISIX-based ingress controller.
- Avi Kubernetes Operator provides L4-L7 load-balancing using VMware NSX Advanced Load Balancer.
- BFE Ingress Controller is a BFE-based ingress controller.
- Cilium Ingress Controller is an ingress controller powered by [Cilium](#).
- The Citrix ingress controller works with Citrix Application Delivery Controller.
- Contour is an Envoy based ingress controller.
- Emissary-Ingress API Gateway is an Envoy-based ingress controller.
- EnRoute is an Envoy based API gateway that can run as an ingress controller.
- Easegress IngressController is an Easegress based API gateway that can run as an ingress controller.
- F5 BIG-IP Container Ingress Services for Kubernetes lets you use an Ingress to configure F5 BIG-IP virtual servers.
- FortiADC Ingress Controller support the Kubernetes Ingress resources and allows you to manage FortiADC objects from Kubernetes
- Gloo is an open-source ingress controller based on Envoy, which offers API gateway functionality.
- HAProxy Ingress is an ingress controller for HAProxy.
- Higress is an Envoy based API gateway that can run as an ingress controller.
- The HAProxy Ingress Controller for Kubernetes is also an ingress controller for HAProxy.
- Istio Ingress is an Istio based ingress controller.
- The Kong Ingress Controller for Kubernetes is an ingress controller driving Kong Gateway.
- Kusk Gateway is an OpenAPI-driven ingress controller based on Envoy.
- The NGINX Ingress Controller for Kubernetes works with the NGINX webserver (as a proxy).
- The ngrok Kubernetes Ingress Controller is an open source controller for adding secure public access to your K8s services using the ngrok platform.
- The OCI Native Ingress Controller is an Ingress controller for Oracle Cloud Infrastructure which allows you to manage the OCI Load Balancer.
- The Pomerium Ingress Controller is based on Pomerium, which offers context-aware access policy.
- Skipper HTTP router and reverse proxy for service composition, including use cases like Kubernetes Ingress, designed as a library to build your custom proxy.
- The Traefik Kubernetes Ingress provider is an ingress controller for the Traefik proxy.
- Tyk Operator extends Ingress with Custom Resources to bring API Management capabilities to Ingress. Tyk Operator works with the Open Source Tyk Gateway & Tyk Cloud control plane.
- Voyager is an ingress controller for HAProxy.
- Wallarm Ingress Controller is an Ingress Controller that provides WAAP (WAF) and API Security capabilities.

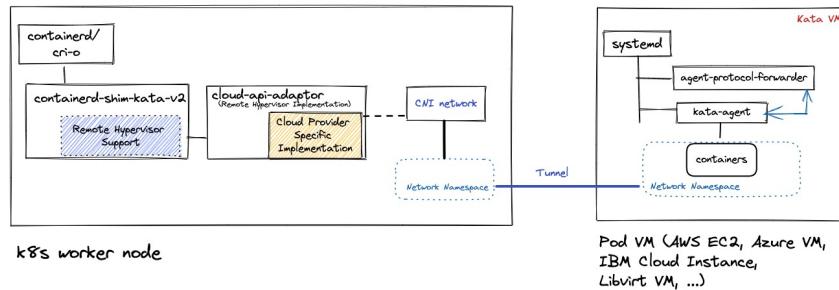
<https://kubernetes.io/docs/concepts/services-networking/ingress-controllers/>

# Lab

- Configuring eBPF With Cilium

# Section 3: Network Troubleshooting

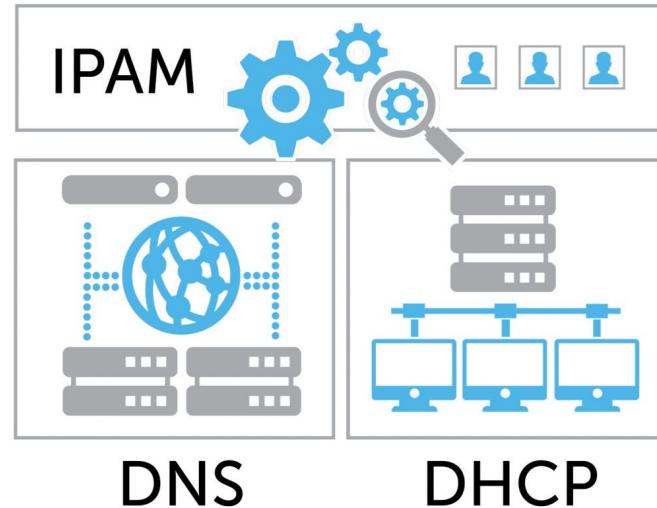
- Host troubleshooting
- Internal k8s network troubleshooting



# IPAM For Kubernetes

- Management and assign of IP addresses

DDI integrates DNS, DHCP, and IPAM into one solution



# DNS Breakdown

- **Service:** my-svc.my-namespace.svc.cluster-domain.example
- **Pod:** pod-ip-address.my-namespace.pod.cluster-domain.example
- **Pod created by a deployment exposed as a service:**  
pod-ip-address.deployment-name.my-namespace.svc.cluster-domain.example

# Host Network Troubleshooting

- Check subnets
- Check private/public IP addresses for Control Plane
- Check firewall rules

# Check Overall Network

- Ensure that the subnets and CIDRs are correct

# Check Private/Public IPs

- Ensure that you're running private or public IPs
- If you're running private, ensure that you have a way to access them

# Check Firewall Rules

- Are the proper firewall rules open?
- Ensure that not ALL firewall rules are open

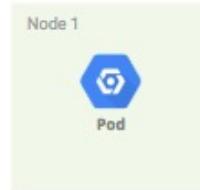
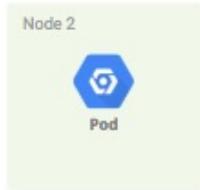
# CNI Troubleshooting

- Check that the Pods for the CNI are running
- Check Pod logs
- Check Pod events

# Are Pods Running?



Node 1 Could Be Out Of Resources, Overflowing To Node 2



Could Be A Daemon Pod, Could Be Luck



All Pods Could Be In One Node If Resources Exist



No Reason Node 1 Is Where We Start Putting Pods

# Pod Port Troubleshooting

- Check which containers are using which Ports
- Ensure that there aren't any Port conflicts

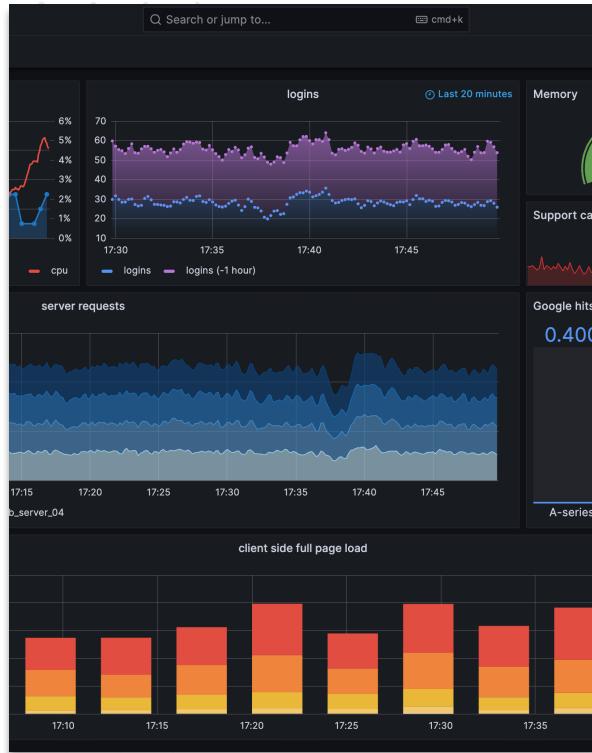
# Pod Troubleshooting

- Describe
- Log
- Top
- Events
- debug

# Lab

- Troubleshooting Pods

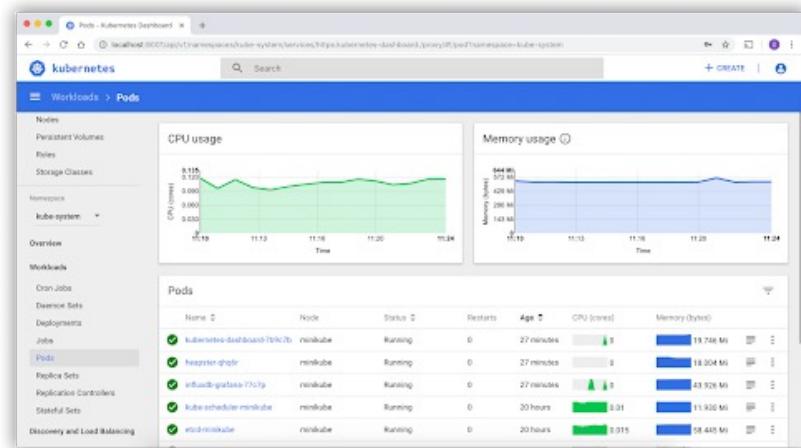
## Section 4: Network Observability



- Why monitoring and observability
- Network observability at the host and Pod level

# Why Monitoring

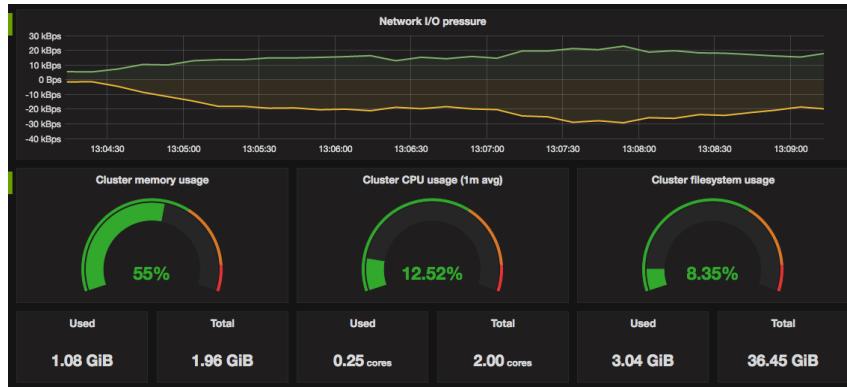
- See what's happening in real-time and take action



# Why Observability

- Automate the action taking

# Monitoring and Observability At The Host Level



- See the overall throughput for the host
- Watch for network blips and any downtime
- High response/MS
- Resource utilization
- VM health

# Network Issues

- Ingress and Egress
- Local VM Networking
- Bandwidth

# Resource Utilization and Scaling

- Proper amount of CPU and Memory available.
- Scalable Worker Nodes

# VM Health

- IF your Worker Nodes aren't Serverless

# Monitoring And Observability At The Pod Level



- Check for Pod networking health
- Ensure No CNI Issues
- Resource Optimization
- Metrics
- Events and Logs

# Pod Networking

- CNI Health

# Resource Optimization

- Proper amount of memory and CPU available for Pods
- Vertical and Horizontal Pod Autoscaling

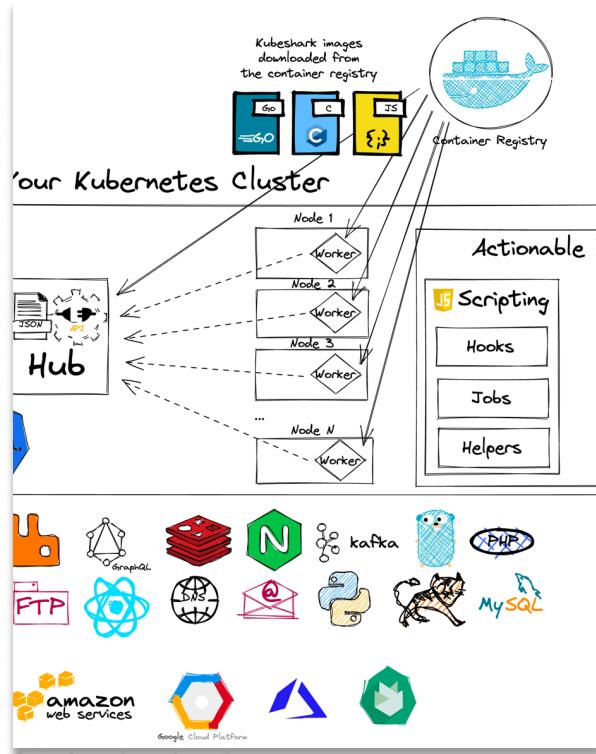
# Events and Logs

- What we already talked about in Section 3:
  - `kubectl logs`
  - `kubectl events`
  - `kubectl top`
  - `kubectl describe`
  - `kubectl debug`

# Lab

- Configuring Grafana
- Setting up resource optimization for Pods

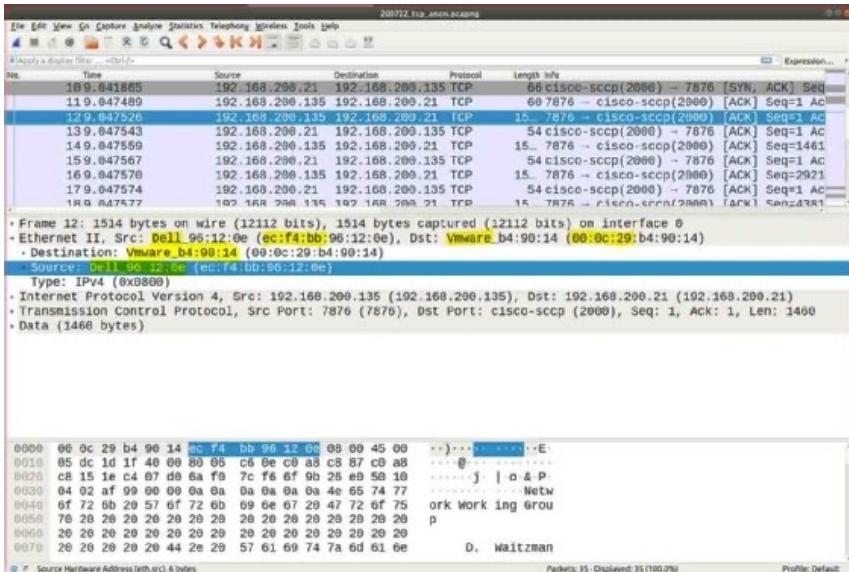
## Section 5: Observability With Kubeshark



- What and why Kubeshark
  - Monitoring and tracing API calls
  - Filtering

# What Is Kubeshark

- Like Wireshark, but for Kubernetes



# Why Use It?

- Tracing API calls
- Monitoring API calls
- Filtering network traffic

# Installing Kubeshark

The screenshot shows the Kubeshark web application running at `localhost:8899/?q=`. The title bar says "Kubeshark" and the address bar shows the URL. A banner at the top right states "Targetting 22 pods in namespaces [kube-system], [sock-shop]".

**Kubeshark Filter Syntax** (top left):

- Streaming live traffic (green icon)
- Apply button
- Copy and Paste icons

**HTTP Requests (left column):**

- 200 GET /health [Unresolved] ↪ payment.sock-shop 172.17.0.1:59994 → 172.17.0.1980
- 200 GET /health [Unresolved] ↪ payment.sock-shop 172.17.0.1:60010 → 172.17.0.1980
- 200 GET /health [Unresolved] ↪ user.sock-shop 172.17.0.1:37452 → 172.17.0.980
- 200 GET /health [Unresolved] ↪ user.sock-shop 172.17.0.1:37466 → 172.17.0.980

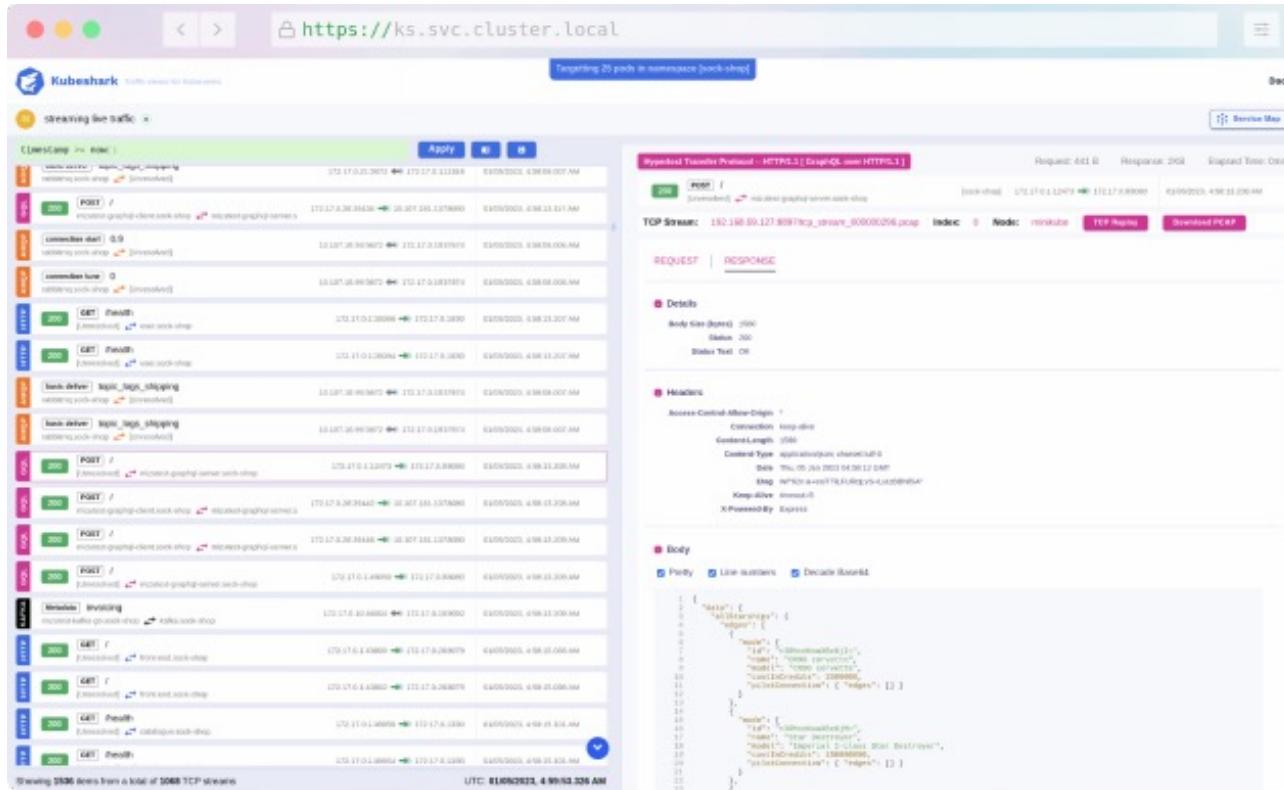
**Hypertext Transfer Protocol -- HTTP/1.1 (right side):**

- 101 GET /ws kubeshark-hub.kubeshark ↪ [Unresolved]
- Stream: 192.168.49.2:8897/tcp\_stream\_000000004

**REQUEST | RESPONSE** (bottom right)

**Details** (bottom right)

# Monitoring API Calls



# Tracing API Calls

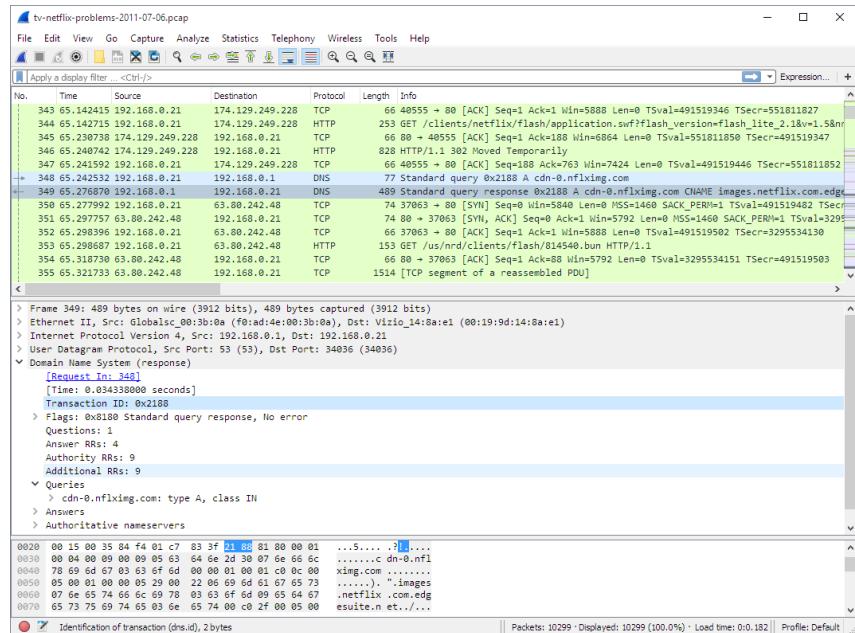
- Having the ability to see who and what is interacting with Kubernetes API endpoints

# Filtering

- Seeing in real-time and past-time down to the second what and who is interacting with your environment

# Wireshark

- Monitor all network traffic

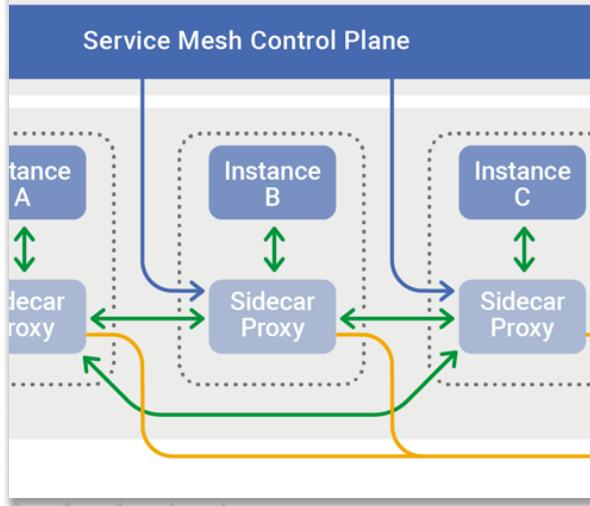


# Lab

- Install and configure Kubeshark
- Deploy a workload and view API calls

## Section 6: Service Mesh

- What's a service mesh
- Service mesh options



# What Is A Service Mesh

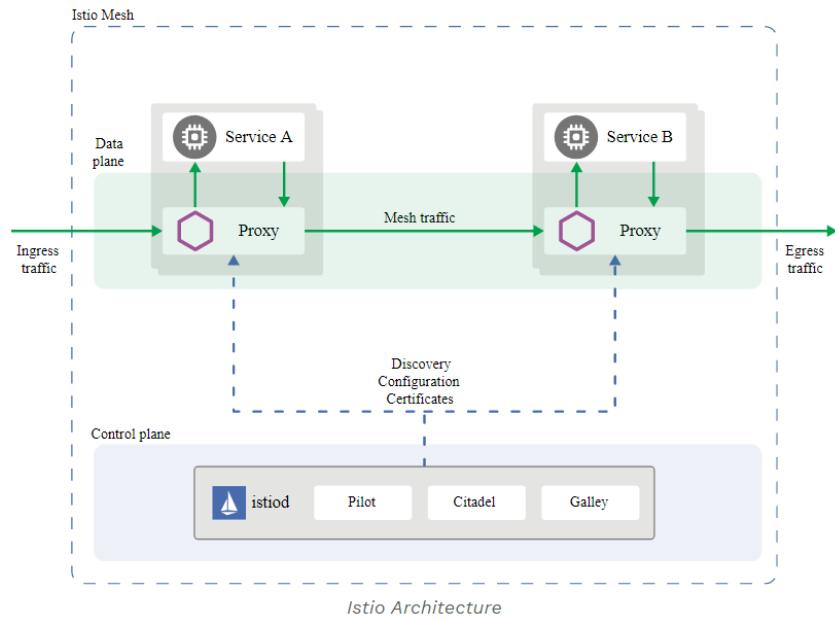
- End-to-end service encryption
- Network traffic management

# Service Mesh Options

- Istio
- Linkerd
- Consul Connect

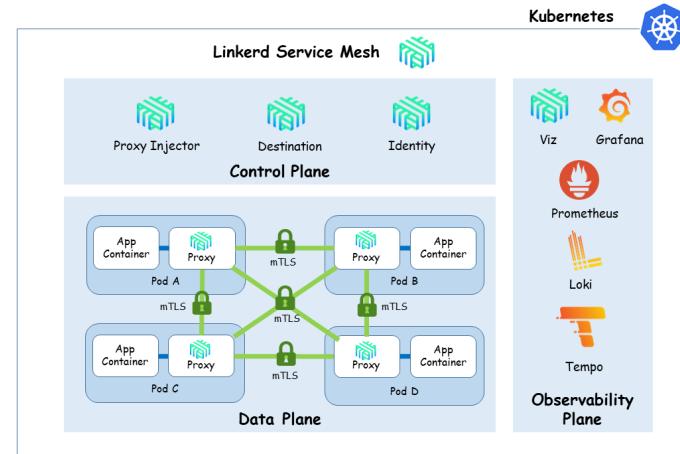
# Istio

- Most popular Service Mesh



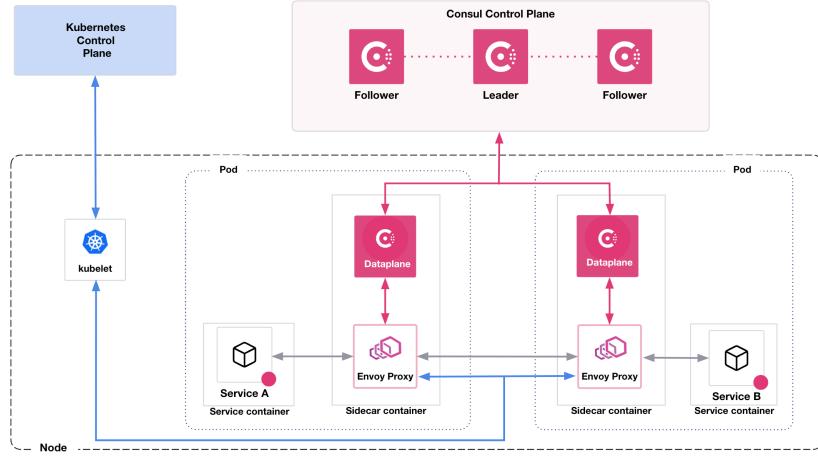
# Linkerd

- Easy to use service mesh  
(service mesh out of the box)



# Consul Connect

- HashiCorps service mesh



# Service Mesh Latency Management

- Manage and capture network latency

# Service Mesh Encryption

- Utilization of mTLS

# Network Policies

- Provide security for ingress/egress
- Block internal and external traffic

# Lab

- Configuring Istio for network latency
- Network Policies